

Statistics for Neuroscience – Lab Activity

Livio

2026-02-20

Activity: Discover how contrasts and centering change your model

Goal

By fitting models with different encodings, *discover* for yourself:

- How dummy vs sum contrasts change coefficient *meaning*
- Why centering continuous predictors matters (especially with interactions)
- How to get “clean” interpretations of main effects

Instructions

Work step-by-step. Compare `summary()` outputs and plots. Explain what you observe in your own words.

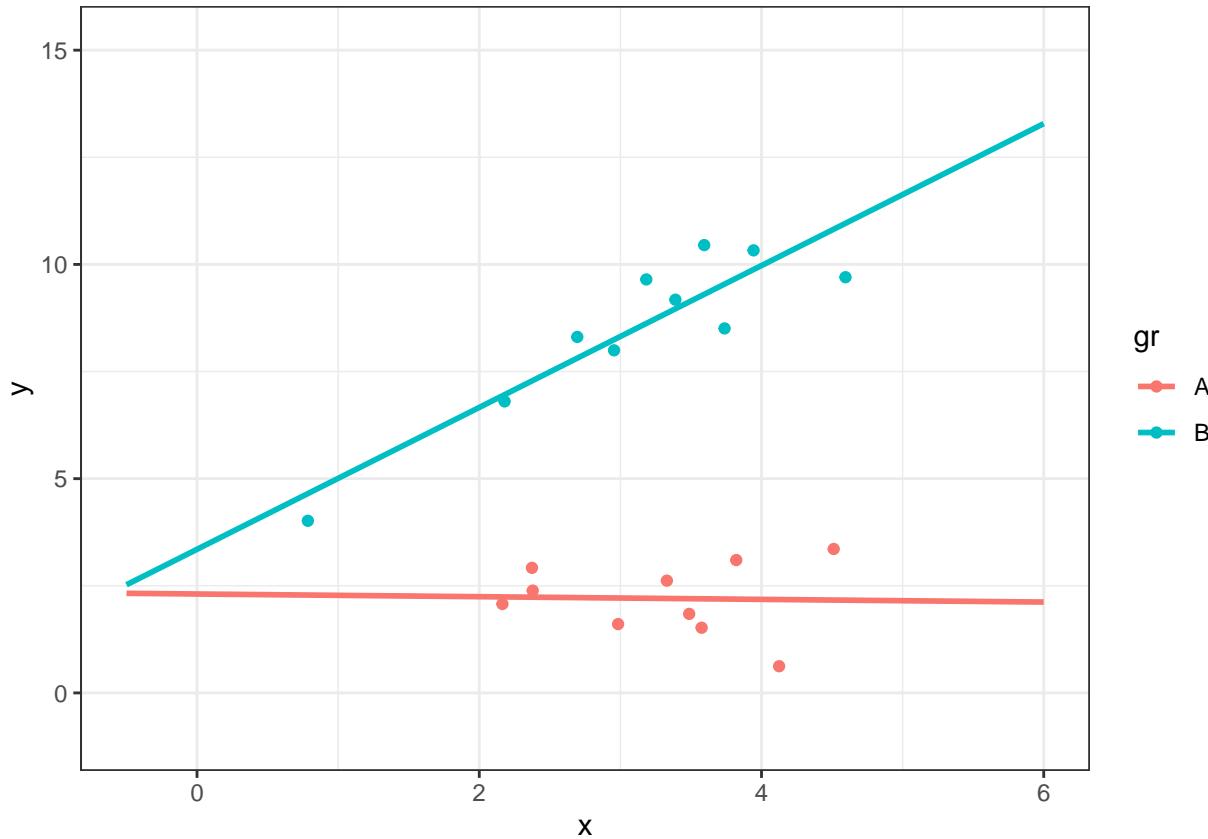
1. Generate the data and visualize it

Run:

```
set.seed(1)
n0 <- 10
D <- data.frame(
  gr = as.factor(rep(LETTERS[1:2], n0)),
  x  = rnorm(n0 * 2) + 3
)
mu <- 2 + (D$gr == "B") * .5 + D$x * (D$gr == "B") * 2
D$y <- mu + rnorm(n0 * 2)

library(ggplot2)
ggplot(D, aes(x = x, y = y, color = gr)) +
  geom_point() +
  geom_smooth(method = "lm", fill = NA, fullrange = TRUE) +
  xlim(-.5, 6) +
  theme_bw()

## `geom_smooth()` using formula = 'y ~ x'
```



Tasks

1. Write down (from the code) the *true* generative model for groups A and B (two separate regression equations).
2. From the plot, describe:
 - Does x affect y in group A?
 - Does x affect y in group B?
 - Is there an interaction between gr and x?

2. Model 1: Default dummy coding (no centering)

```
mod1 <- lm(y ~ gr * x, data = D)
summary(mod1)

##
## Call:
## lm(formula = y ~ gr * x, data = D)
##
## Residuals:
##      Min       1Q   Median       3Q      Max 
## -1.55585 -0.61528 -0.00131  0.54234  1.19203 
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 2.30818   1.24368   1.856  0.08198 .  
## grB          0.22000   0.11200   1.964  0.05482 .  
## x            1.00000   0.11200  8.929  <2e-16 ***
## grB:x       0.77000   0.15600  4.963  0.00018 ***
```

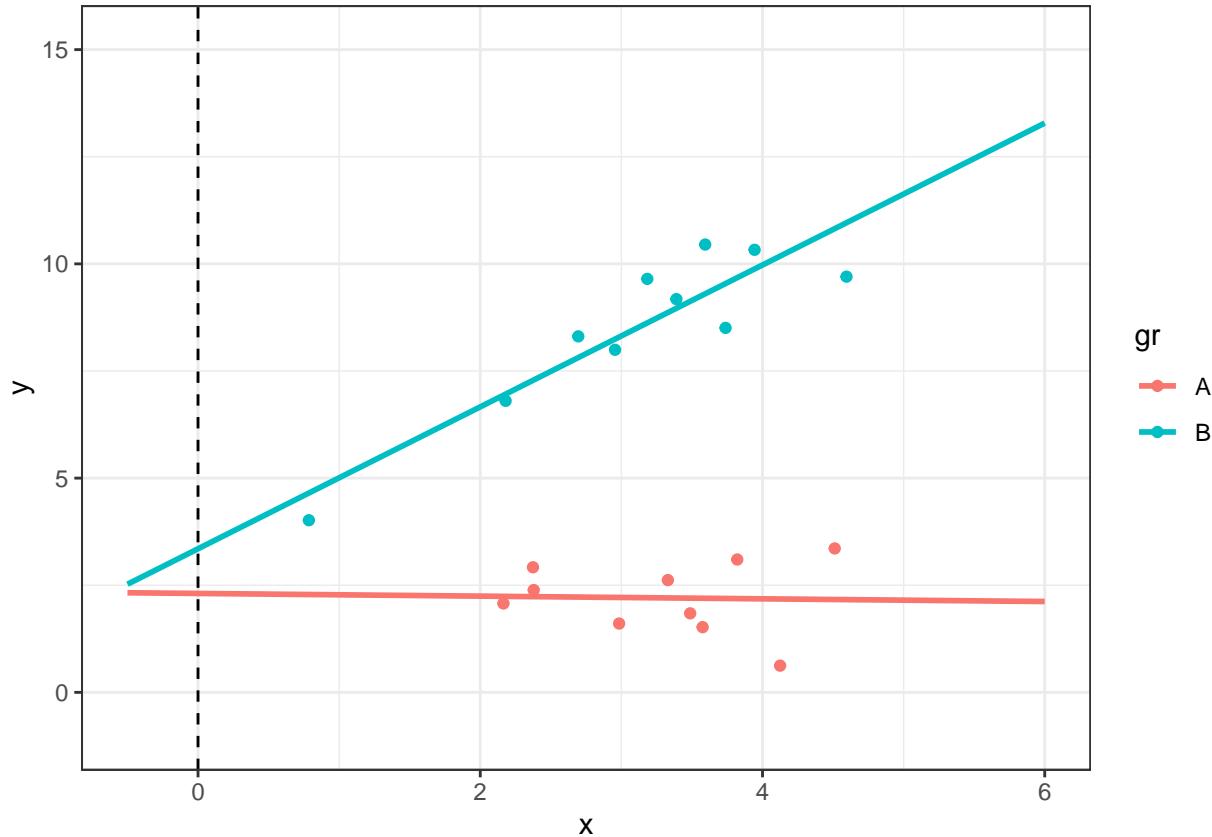
```

## grB      1.04292   1.53659   0.679   0.50701
## x       -0.03137   0.37016  -0.085   0.93352
## grB:x    1.68703   0.46200   3.652   0.00215 ** 
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.8778 on 16 degrees of freedom
## Multiple R-squared:  0.9481, Adjusted R-squared:  0.9384
## F-statistic: 97.49 on 3 and 16 DF,  p-value: 1.708e-10

p2 <- ggplot(D, aes(x = x, y = y, color = gr)) +
  geom_point() +
  geom_smooth(method = "lm", fill = NA, fullrange = TRUE) +
  geom_vline(xintercept = 0, linetype = "dashed") +
  xlim(-.5, 6) +
  theme_bw()
print(p2)

```

`geom_smooth()` using formula = 'y ~ x'



Discover

1. What does the grB coefficient mean? (Hint: at what x?)
2. Is the group difference meaningful given where data are?
3. Look at p-values. Which effects are “significant”? Interaction?

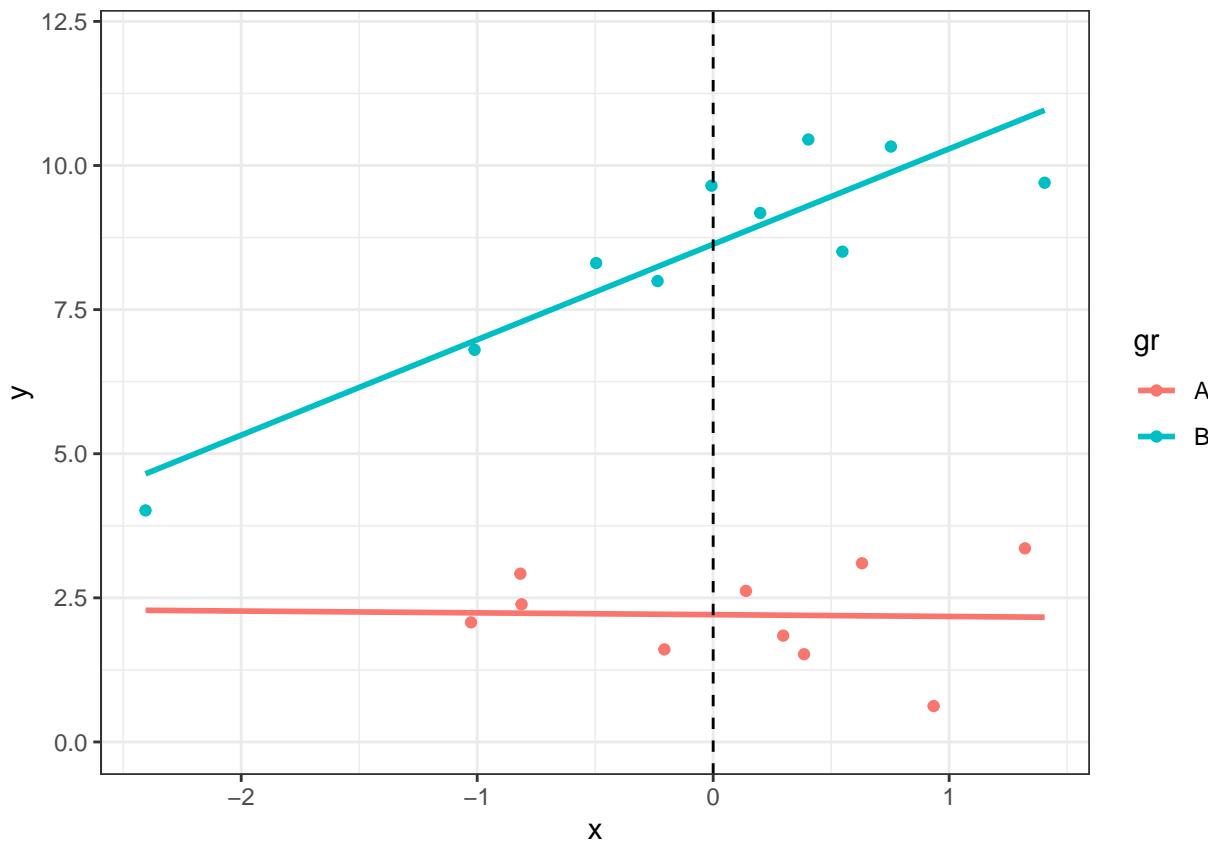
3. Model 2: Center x (same dummy coding)

```
D2 <- D
D2$x <- D$x - mean(D$x)
mod2 <- lm(y ~ gr * x, data = D2)
summary(mod2)

##
## Call:
## lm(formula = y ~ gr * x, data = D2)
##
## Residuals:
##       Min     1Q Median     3Q    Max 
## -1.55585 -0.61528 -0.00131  0.54234  1.19203 
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 2.20810   0.27933   7.905 6.47e-07 ***
## grB          6.42543   0.39449  16.288 2.21e-11 ***
## x            -0.03137   0.37016  -0.085  0.93352    
## grB:x        1.68703   0.46200   3.652  0.00215 **  
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.8778 on 16 degrees of freedom
## Multiple R-squared:  0.9481, Adjusted R-squared:  0.9384 
## F-statistic: 97.49 on 3 and 16 DF,  p-value: 1.708e-10

p3 <- ggplot(D2, aes(x = x, y = y, color = gr)) +
  geom_point() +
  geom_smooth(method = "lm", fill = NA, fullrange = TRUE) +
  geom_vline(xintercept = 0, linetype = "dashed") +
  theme_bw()
print(p3)

## `geom_smooth()` using formula = 'y ~ x'
```



Discover

1. Now what does grB mean? (Compare to plot.) Better?
 2. How do p-values change for main effects? For interaction?
 3. Why did centering x change the *main effect* tests, even though interaction was always in the model?
-

4. Model 3: Zero-sum contrasts for gr (+ centered x)

```
D3 <- D2
contrasts(D3$gr) <- contr.sum(2)
mod3 <- lm(y ~ gr * x, data = D3)
summary(mod3)

##
## Call:
## lm(formula = y ~ gr * x, data = D3)
##
## Residuals:
##      Min       1Q   Median       3Q      Max 
## -1.55585 -0.61528 -0.00131  0.54234  1.19203 
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept)  5.4208    0.1972  27.483 6.80e-15 ***
## gr1         -3.2127    0.1972 -16.288 2.21e-11 ***
## x            0.8121    0.2310   3.516  0.00287 ** 
## gr1:x       -0.8435    0.2310  -3.652  0.00215 **
```

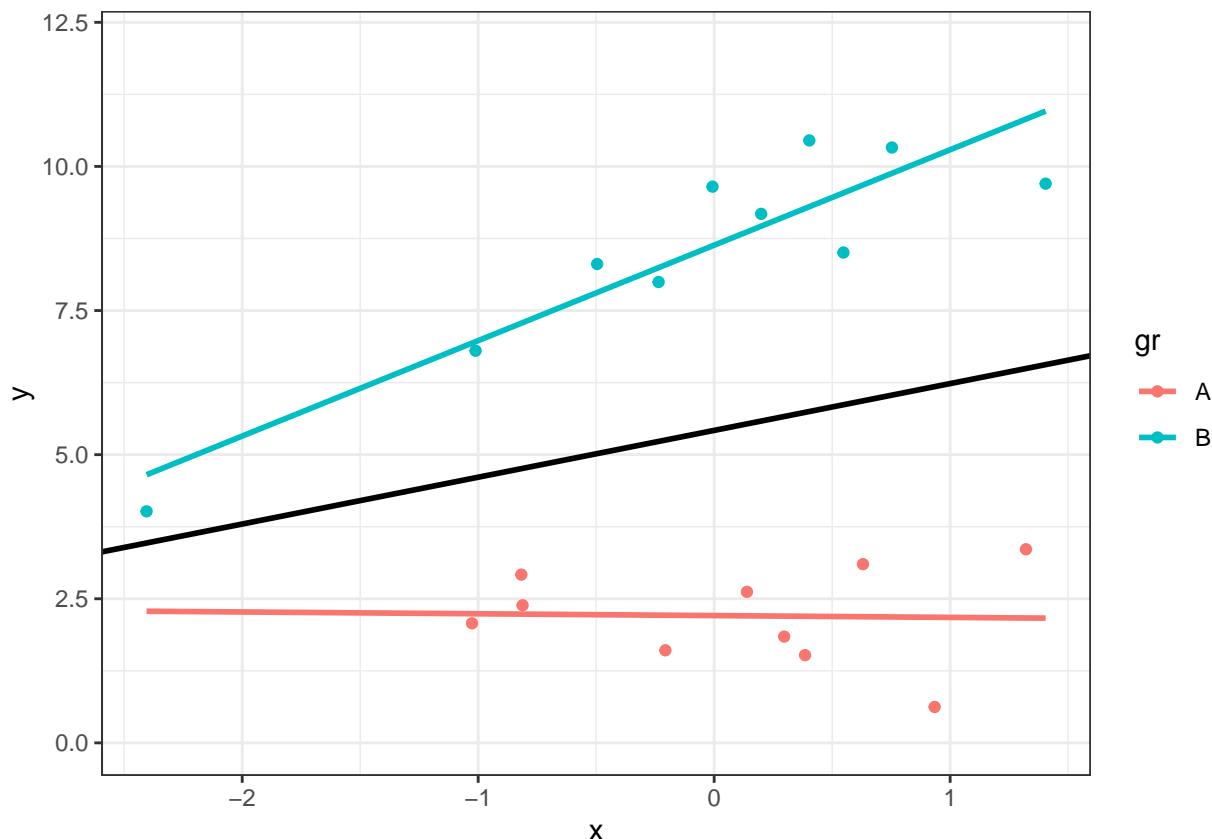
```

## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.8778 on 16 degrees of freedom
## Multiple R-squared:  0.9481, Adjusted R-squared:  0.9384
## F-statistic: 97.49 on 3 and 16 DF,  p-value: 1.708e-10

p4 <- ggplot(D3, aes(x = x, y = y, color = gr)) +
  geom_point() +
  geom_smooth(method = "lm", fill = NA, fullrange = TRUE) +
  geom_abline(intercept = coef(mod3)[1], slope = coef(mod3)[3], color = "black", linewidth = 1) +
  theme_bw()
print(p4)

```

`geom_smooth()` using formula = 'y ~ x'



Discover

1. What population does the intercept + x slope (black line) describe?
2. How is this different from Models 1–2?
3. Compare p-values across all 3 models. Which encoding gives “cleanest” tests?

5. Bonus: Rescale x instead of just centering

```

D4 <- D
D4$x <- scale(D$x) # Centers AND divides by SD
mod4 <- lm(y ~ gr * x, data = D4)
summary(mod4)

```

```

## 
## Call:
## lm(formula = y ~ gr * x, data = D4)
##
## Residuals:
##    Min     1Q Median     3Q    Max
## -1.55585 -0.61528 -0.00131  0.54234  1.19203
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 2.20810   0.27933   7.905 6.47e-07 ***
## grB         6.42543   0.39449  16.288 2.21e-11 ***
## x          -0.02865   0.33805  -0.085  0.93352
## grB:x       1.54069   0.42193   3.652  0.00215 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.8778 on 16 degrees of freedom
## Multiple R-squared:  0.9481, Adjusted R-squared:  0.9384
## F-statistic: 97.49 on 3 and 16 DF,  p-value: 1.708e-10

```

Discover

1. Compare `mod4` coefficients to `mod2` (just centering). Why are slopes $\sim 2x$ larger now?
 2. Is the group difference (`grB`) at $x=0$ the *same*?
 3. **Pro:** What does “per 1-SD change in x ” mean for interpretation?
 4. **Con:** Does rescaling change p-values or tests? Why/why not? When would you rescale?
-

6. Your discoveries

Answer these (no peeking at solutions):

1. Why is the “main effect” of `gr` meaningless without specifying *at what x?*
2. How does centering `x` make main effects interpretable **despite** interaction?
3. When interactions are present, why do defaults hurt power/interpretation?
4. Rule of thumb: when to center? When to use `contr.sum()`?

Extension (groups): Fit on *your own* simulated data (change slopes/groups). Does the lesson hold?

Answer to your side question: Rescaling vs centering

Rescaling (`scale(x)` = centering + dividing by SD) has these effects:

Pros:

- Coefficients are **standardized effect sizes** (“per 1-SD change in x ”). Great for comparing predictor strength (e.g., BMI vs age).
- Group difference still at meaningful $x=0$ (mean).

Cons vs pure centering:

- **No change to p-values, F-tests, or R²** (pure linear rescaling). Tests identical to centered-only model.
- Slopes scale up by $1/\text{SD}_x$, so intercept adjusts down accordingly.
- Less intuitive if SD lacks meaning (e.g., arbitrary imaging parameter).

When to rescale: For predictor comparison or publication tables. Otherwise, centering suffices for clean main effects.